AI Safety, Politics and Policy

Course Number TBD

Instructor: Adam Bouyamourn
Center for Data Science, New York University

Spring 2027

Course description

AI is a fast-moving technical field with broad social implications. Understanding the social implications of AI technologies involves grounding in both the engineering aspects of LLM systems, and the broader theoretical frameworks used to state and understand social problems, as well considering how these problems are concretely addressed by political institutions and actors. This course considers the demand for and supply of AI regulation. We consider, first, the social problems that AI may or may not affect. We then consider engineering approaches used by frontier labs in their attempt to build AI that is "harmless and helpful". Finally, we consider AI policymaking in practice: identifying the key decision-makers in different regulatory domains, understanding their incentives, and evaluating their track records in regulating emerging technologies.

Course schedule

Week	Topic
Part I. Problems: The Demand for Regulation	
1	The Alignment Problem as a Policy Problem: Incomplete contracting, principal-agent problems, and the demand for and supply of regulation
2	If Anyone Builds It, Does Everyone Die?
3	AI, automation, and the macroeconomy
4	AI and democracy: Persuasion, manipulation, disinformation
5	Race, NLP, and fairness in automated decision systems
Part II. Engineering: Frontier Labs and Self-Regulation	
6	What is an LLM? Transformers, reasoning, and chain-of-thought monitorability
7	How LLMs are steered: RLHF, RLAIF, Constitutional AI, activation steering
8	Steering methods and their failures: sycophancy, jailbreaks, abliteration
9	Evaluation and audits
Part III. Regulation in Practice: the Supply of Regulation	
10	Tech, antitrust, and the market structure of AI
11	Who governs? Agenda-setting, GDPR, the EU AI Act, California
12	National security I: drones, warfare, surveillance
13	National security II: chips, compute, China, model diffusion
Part IV. Synthesis	
14	The regulatory alignment problem

Assignments and grading

Logistics. TBD.

Assignments. Students are responsible for three short response memos and one longer final essay. Each

memo (˜1250 words) responds to the required readings and discussion items for one week of the student's choice; the memos must cover different weeks. The final essay (˜3,000 words) is drawn from a pool of overview questions. Students are permitted to suggest their own overview question, but are gently discouraged from doing so. For both the response memos and the essay, you are expected to engage with the optional reading from the relevant weeks.

Grading.

- Three response memos (any two weeks of the student's choice): 60%
- Final essay engaging the readings from any two weeks: 40%

PART I. PROBLEMS — THE DEMAND FOR REGULATION

Week 1. Foundations: principal–agent, capture, and the economic theory of regulation

What tools from public policy can we use to understand the challenges of AI? What new challenges does AI pose? What are the big picture risks and open questions? How should we even think about AI safety?

Required

- G. J. Stigler (1971). "The Theory of Economic Regulation." In: *Bell Journal of Economics and Management Science* 2.1, pp. 3–21. DOI: 10.2307/3003160.
- B. Holmström (1979). "Moral Hazard and Observability." In: *Bell Journal of Economics* 10.1, pp. 74–91. DOI: 10.2307/3003320.
- D. Hadfield-Menell and G. K. Hadfield (2019). "Incomplete Contracting and AI Alignment." In: *Proceedings of the 2019 AAAI/ACM Conference on AI, Ethics, and Society (AIES)*. ACM, pp. 417–422. arXiv: 1804.04268.
- N. Collina et al. (2025). *Emergent Alignment via Competition*. arXiv preprint. v2, February 2026. arXiv: 2509.15090.
- B. Christian (2020). *The Alignment Problem: Machine Learning and Human Values*. New York: W. W. Norton. URL: https://wwnorton.com/books/9780393635829, Chapter 5 "Shaping".

Optional

- S. Peltzman (1976). "Toward a More General Theory of Regulation." In: *Journal of Law and Economics* 19.2, pp. 211–240. DOI: 10.1086/466865.
- D. Carpenter (2014). "Detecting and Measuring Capture." In: *Preventing Regulatory Capture: Special Interest Influence and How to Limit It*. Ed. by D. Carpenter and D. A. Moss. New York: Cambridge University Press, pp. 57–68. DOI: 10.1017/CBO9781139565875.006 (chapter in Carpenter and Moss, 2014).

Discussion

Is AI alignment a new problem? Discuss.

Week 2. *If Anyone Builds It, Does Everyone Die?*

Doom, X-Risk, and AI-induced catastrophes. What picture of the world, and of AI capabilities, do you need in order to believe in catastrophic AI risk? Is such a picture adequate?

Required

- A. Narayanan and S. Kapoor (2025). "AI as Normal Technology." In: *Knight First Amendment Institute*. URL: https://knightcolumbia.org/content/ai-as-normal-technology.

- D. Hendrycks, M. Mazeika, and T. Woodside (2023). *An Overview of Catastrophic AI Risks*. arXiv preprint. arXiv: 2306.12001.
- J. Carlsmith (2022). *Is Power-Seeking AI an Existential Risk?* arXiv preprint. arXiv: 2206.13353.
- E. Yudkowsky and N. Soares (2025). *If Anyone Builds It, Everyone Dies*. Little, Brown. URL: https://en.wikipedia.org/wiki/If_Anyone_Builds_It,_Everyone_Dies
- M. Mitchell (2021). "Why AI Is Harder than We Think." In: *arXiv preprint*. arXiv: 2104.12871.

Optional

- A. Narayanan and S. Kapoor (2024). *AI Snake Oil: What Artificial Intelligence Can Do, What It Can't, and How to Tell the Difference*. Princeton: Princeton University Press. URL: https://press.princeton.edu/books/hardcover/9780691249131/ai-snake-oil, relevant chapter.
- Review of G. Lovely (2026). *Obsolete: The AI Industry's Trillion Dollar Race to Replace You — And How to Stop It*. Forthcoming, September 2026. OR Books. URL: https://www.obsoletebook.org/ (*TBD; book due Sept 2026*).
- H. G. Frankfurt (2005). *On Bullshit*. Princeton: Princeton University Press. URL: https://press.princeton.edu/books/hardcover/9780691122946/on-bullshit.

Discussion

If anyone builds it, does everyone die?

Week 3. AI, automation, and the macroeconomy

How do economic capabilities differ from technological capabilites? Is AI coming for our jobs? If so, whose? What can we learn from economic history about technological disruptions? What *can't* AI do?

Required

- P. A. David (1990). "The Dynamo and the Computer: An Historical Perspective on the Modern Productivity Paradox." In: *American Economic Review (Papers and Proceedings)* 80.2, pp. 355–361. DOI: 10.2307/2006600.
- T. F. Bresnahan and M. Trajtenberg (1995). "General Purpose Technologies: "Engines of Growth"?" In: *Journal of Econometrics* 65.1, pp. 83–108. DOI: 10.1016/0304-4076(94)01598-T.
- D. Acemoglu and P. Restrepo (2020). "Robots and Jobs: Evidence from US Labor Markets." In: *Journal of Political Economy* 128.6, pp. 2188–2244. DOI: 10.1086/705716.
- D. Acemoglu and S. Johnson (2024). "Learning From Ricardo and Thompson: Machinery and Labor in the Early Industrial Revolution and in the Age of Artificial Intelligence." In: *Annual Review of Economics* 16.Volume 16, 2024, pp. 597–621. DOI: https://doi.org/10.1146/annurev-economics-091823-025129. URL: https://www.annualreviews.org/content/journals/10.1146/annurev-economics-091823-025129

Optional

- N. Crafts (2004). "Steam as a General Purpose Technology: A Growth Accounting Perspective." In: *Economic Journal* 114.495, pp. 338–351. DOI: 10.1111/j.1468-0297.2004.00200.x.
- D. Acemoglu (2024). *The Simple Macroeconomics of AI*. NBER Working Paper 32487. National Bureau of Economic Research. URL: https://www.nber.org/papers/w32487.
- E. Brynjolfsson, D. Rock, and C. Syverson (2021). "The Productivity J-Curve: How Intangibles Complement General Purpose Technologies." In: *American Economic Journal: Macroeconomics* 13.1, pp. 333–372. DOI: 10.1257/mac.20180386.
- A. Korinek and J. E. Stiglitz (2017). *Artificial Intelligence and Its Implications for Income Distribution and Unemployment*. NBER Working Paper 24174. National Bureau of Economic Research. URL: https://www.nber.org/papers/w24174.
- R. W. Fogel (1964). *Railroads and American Economic Growth: Essays in Econometric History*. Baltimore: Johns Hopkins University Press. URL: https://eh.net/book_reviews/railroads-and-american-economic-growth-essays-in-econometric-history/, summary via review.

Discussion

How will AI affect the labor market, technological innovation, and economic growth? Is AI different to previous economic technologies? Why or why not?

Week 4. AI and Democracy: Persuasion, manipulation, disinformation

How does AI affect public discourse? Are we entering a new age of disinformation? Alternatively, can AI be used to improve democracy?

Required

- H. Waight et al. (2026). "State Media Control Influences Large Language Models." In: *Nature*. DOI: 10.1038/s41586-026-10506-7.
- J. A. Goldstein et al. (2023). *Generative Language Models and Automated Influence Operations: Emerging Threats and Potential Mitigations*. CSET Report, Center for Security and Emerging Technology, Georgetown University. arXiv: 2301.04246.
- M. Burtell and T. Woodside (2023). *Artificial Influence: An Analysis of AI-Driven Persuasion*. arXiv preprint. arXiv: 2303.08721.
- C. Summerfield et al. (2025). "The Impact of Advanced AI Systems on Democracy." In: *Nature Human Behaviour* 9, pp. 2420–2430. DOI: 10.1038/s41562-025-02309-z.

Optional

- D. M. J. Lazer et al. (2018). "The Science of Fake News." In: *Science* 359.6380, pp. 1094–1096. DOI: 10.1126/science.aao2998.
- J. Hartmann, J. Schwenzow, and M. Witte (2023). *The Political Ideology of Conversational AI: Converging Evidence on ChatGPT's Pro-environmental, Left-libertarian Orientation*. arXiv preprint. arXiv: 2301.01768.

Discussion

Will AI harm or improve public discourse? Does it allow for new direct democratic possibilities?

Week 5. Race, NLP, and fairness in automated decision systems

LLM pipelines are already being used to make high-stakes decisions. Can we trust them?

Required

- S. L. Blodgett et al. (2020). "Language (Technology) Is Power: A Critical Survey of "Bias" in NLP." in: *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)*, pp. 5454–5476. URL: https://aclanthology.org/2020.acl-main.485/.
- M. Coots et al. (2025). "Racial Bias in Clinical and Population Health Algorithms: A Critical Review of Current Debates." In: *Annual Review of Public Health* 46, pp. 507–523. DOI: 10.1146/annurev-publhealth-071823-112058.
- J. D. Gaebler et al. (2024). "Auditing Large Language Models for Race & Gender Disparities: Implications for Artificial Intelligence-Based Hiring." In: *Socius: Sociological Research for a Dynamic World*. DOI: 10.1177/23794607251320229.
- M. Hardt, E. Price, and N. Srebro (2016). "Equality of Opportunity in Supervised Learning." In: *Advances in Neural Information Processing Systems 29 (NeurIPS)*, pp. 3315–3323. arXiv: 1610.02413.

Optional

- L. T. Liu, I. D. Raji, A. Zhou, et al. (2026). *Bridging Prediction and Intervention Problems in Social Systems*. arXiv preprint. Workshop position paper from the Simons/BIRS workshop on Bridging Prediction and Intervention Problems in Social Systems; 38 authors. arXiv: 2507.05216.
- A. Chouldechova (2017). "Fair Prediction with Disparate Impact: A Study of Bias in Recidivism Prediction Instruments." In: *Big Data* 5.2, pp. 153–163. DOI: 10.1089/big.2016.0047.
- J. Kleinberg, S. Mullainathan, and M. Raghavan (2017). "Inherent Trade-Offs in the Fair Determination of Risk Scores." In: *Proceedings of the 8th Innovations in Theoretical Computer Science Conference (ITCS)*. arXiv: 1609.05807.
- M. Sap et al. (2019). "The Risk of Racial Bias in Hate Speech Detection." In: *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics (ACL)*, pp. 1668–1678. URL: https://aclanthology.org/P19-1163/.

- H. H. Nigatu and I. D. Raji (2024). "I Searched for a Religious Song in Amharic and Got Sexual Content Instead: Investigating Online Harm in Low-Resourced Languages on YouTube." In: *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency (FAccT)*. DOI: 10.1145/3630106.3658546.
- Z. Obermeyer et al. (2019). "Dissecting Racial Bias in an Algorithm Used to Manage the Health of Populations." In: *Science* 366.6464, pp. 447–453. DOI: 10.1126/science.aax2342.
- L. Blattner, S. Nelson, and J. Spiess (2024). *Unpacking the Black Box: Regulating Algorithmic Decisions*. arXiv preprint. Revised May 2024. arXiv: 2110.03443.
- V. Eubanks (2018). *Automating Inequality: How High-Tech Tools Profile, Police, and Punish the Poor*. New York: St. Martin's Press. URL: https://virginia-eubanks.com/automating-inequality/, relevant chapter.

Discussion

What does the fairness literature teach us about making fair decisions? Can LLMs be used to make fair automated decisions at scale?

Part II. Engineering — self-regulation from frontier labs

Week 6. What is an LLM? Transformers, reasoning, and chain-of-thought monitorability

How did we get here? LLMs are a revolutionary technology: and AI is a cluster term for a suite of methods that are wrapped around an LLM, and a handful of related technologies. How do they work?

Required

- A. Vaswani et al. (2017). "Attention Is All You Need." In: *Advances in Neural Information Processing Systems 30 (NeurIPS)*, pp. 5998–6008. arXiv: 1706.03762.
- L. Bereska and E. Gavves (2024). *Mechanistic Interpretability for AI Safety — A Review*. arXiv preprint. arXiv: 2404.14082.
- T. Korbak et al. (2025). *Chain of Thought Monitorability: A New and Fragile Opportunity for AI Safety*. arXiv preprint. 41-author multi-laboratory position paper. arXiv: 2507.11473.
- J. Wei et al. (2022). "Emergent Abilities of Large Language Models." In: *Transactions on Machine Learning Research*. arXiv: 2206.07682,
- R. Schaeffer, B. Miranda, and S. Koyejo (2023). "Are Emergent Abilities of Large Language Models a Mirage?" In: *Advances in Neural Information Processing Systems 36 (NeurIPS)*. arXiv: 2304.15004.

Optional

- E. M. Bender and A. Koller (2020). "Climbing Towards NLU: On Meaning, Form, and Understanding in the Age of Data." In: *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)*, pp. 5185–5198. URL: https://aclanthology.org/2020.acl-main.463/.
- J. Chen et al. (2026). *Reasoning Models Struggle to Control their Chains of Thought*. arXiv preprint. arXiv: 2603.05706.
- C. R. Shalizi (2026). *"Attention", "Transformers", in Neural Network "Large Language Models"*. Notebook, bactra.org. Last update 19 May 2026; first version March 2023. URL: https://bactra.org/notebooks/nn-attention-and-transformers.html. Sections "Attention" (through Identification Failures), "Transformers", and "It's Just a Markov Model".
- E. M. Bender, T. Gebru, et al. (2021). "On the Dangers of Stochastic Parrots: Can Language Models Be Too Big?" In: *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency (FAccT)*. ACM, pp. 610–623. DOI: 10.1145/3442188.3445922.

Discussion

How do LLMs reason? Do we know? What are the implications for safety?

Week 7. How LLMs are steered: RLHF, RLAIF, Constitutional AI, activation steering

Required

- L. Ouyang et al. (2022). "Training Language Models to Follow Instructions with Human Feedback." In: *Advances in Neural Information Processing Systems 35 (NeurIPS)*. arXiv: 2203.02155.
- S. Casper et al. (2023). "Open Problems and Fundamental Limitations of Reinforcement Learning from Human Feedback." In: *Transactions on Machine Learning Research*. arXiv: 2307.15217.
- Y. Bai et al. (2022). *Constitutional AI: Harmlessness from AI Feedback*. Anthropic technical report. arXiv: 2212.08073.
- V. Conitzer et al. (2024). *Social Choice Should Guide AI Alignment in Dealing with Diverse Human Feedback*. arXiv preprint. arXiv: 2404.10271.
- A. M. Turner et al. (2023). *Steering Language Models with Activation Engineering*. arXiv preprint. arXiv: 2308.10248.

Optional

- Anthropic Claude system card (*most recent version*).
- DeepSeek-R1 technical report (*TBD*).

Discussion

What are the main techniques frontier labs use to build safe AI methods?

Week 8. Steering methods and their failures: sycophancy, jailbreaks, abliteration

Required

- W. V. O. Quine (1960). *Word and Object*. Cambridge, MA: MIT Press. URL: https://mitpress.mit.edu/9780262518314/word-and-object/, Chapter I §§1–2 ("The Beginnings of Objectivity").
- M. Sharma et al. (2023). *Towards Understanding Sycophancy in Language Models*. arXiv preprint. Published at ICLR 2024. arXiv: 2310.13548.
- I. Shapira, G. Benade, and A. D. Procaccia (2026). *How RLHF Amplifies Sycophancy*. Preprint. 31 January 2026. URL: https://procaccia.info/wp-content/uploads/2026/01/sycophancy.pdf
- A. Arditi et al. (2024). *Refusal in Language Models Is Mediated by a Single Direction*. arXiv preprint. arXiv: 2406.11717.
- P. Wiese (2024). *Heretic: A Production Tool for Activation-Based Refusal Removal in Language Models*. Software repository. URL: https://github.com/p-e-w/heretic.

Optional

- L. Gao, J. Schulman, and J. Hilton (2023). "Scaling Laws for Reward Model Overoptimization." In: *Proceedings of the 40th International Conference on Machine Learning (ICML)*. arXiv: 2210.10760.
- J. Wen et al. (2025). "Language Models Learn to Mislead Humans via RLHF." in: *International Conference on Learning Representations (ICLR)*. arXiv: 2409.12822.
- A. Wei, N. Haghtalab, and J. Steinhardt (2023). "Jailbroken: How Does LLM Safety Training Fail?" In: *Advances in Neural Information Processing Systems 36 (NeurIPS)*. arXiv: 2307.02483.
- E. Perez et al. (2023). "Discovering Language Model Behaviors with Model-Written Evaluations." In: *Findings of the Association for Computational Linguistics (ACL)*. arXiv: 2212.09251.

Discussion

Are engineering approaches to AI safety adequate? What can and can't they achieve?

Week 9. Evaluation and audits

It is not straightforward to describe what LLMs can do. It is also not straightforward to measure or benchmark LLM capabilities.

Required

- METR (Model Evaluation and Threat Research) (2025). *Measuring the Ability of AI Agents to Complete Long Tasks*. METR research report. URL: https://metr.org/blog/2025-03-19-measuring-ai-ability-to-complete-long-tasks/.
- METR (Model Evaluation and Threat Research) (2024). *Evaluating Language-Model Agents on Realistic Autonomous Tasks*. METR research report. URL: https://metr.org/research/.
- I. D. Raji, E. M. Bender, et al. (2021). "AI and the Everything in the Whole Wide World Benchmark." In: *Advances in Neural Information Processing Systems 34, Datasets and Benchmarks Track (NeurIPS D&B)*. arXiv: 2111.15366.
- A. Meinke et al. (2024). *Frontier Models Are Capable of In-context Scheming*. Apollo Research technical report. arXiv: 2412.04984.
- A. Alaa et al. (2025). "Medical Large Language Model Benchmarks Should Prioritize Construct Validity." In: *Proceedings of the 42nd International Conference on Machine Learning (ICML)*. Oral presentation; final author list to be confirmed. URL: https://icml.cc/virtual/2025/papers.html.

Optional

- Apollo Research and OpenAI (2025). *Detecting and Reducing Scheming in AI Models*. Joint Apollo Research and OpenAI report, September 2025. URL: https://openai.com/index/detecting-and-reducing-scheming-in-ai-models/.
- I. D. Raji, I. E. Kumar, et al. (2022). "The Fallacy of AI Functionality." In: *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency (FAccT)*, pp. 959–972. DOI: 10.1145/3531146.3533158.
- T. Liao et al. (2021). "Are We Learning Yet? A Meta Review of Evaluation Failures Across Machine Learning." In: *Advances in Neural Information Processing Systems 34, Datasets and Benchmarks Track (NeurIPS D&B)*. URL: https://datasets-benchmarks-proceedings.neurips.cc/paper/2021/hash/757b505cfd34c64c85ca5b5690ee5293-Abstract-round2.html.
- M. Mitchell et al. (2019). "Model Cards for Model Reporting." In: *Proceedings of the 2019 ACM Conference on Fairness, Accountability, and Transparency (FAccT)*, pp. 220–229. DOI: 10.1145/3287560.3287596.
- I. D. Raji, P. Xu, et al. (2022). "Outsider Oversight: Designing a Third Party Audit Ecosystem for AI Governance." In: *Proceedings of the 2022 AAAI/ACM Conference on AI, Ethics, and Society (AIES)*. DOI: 10.1145/3514094.3534181.
- S. Costanza-Chock, I. D. Raji, and J. Buolamwini (2022). "Who Audits the Auditors? Recommendations from a Field Scan of the Algorithmic Auditing Ecosystem." In: *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency (FAccT)*. DOI: 10.1145/3531146.3533213.
- J. B. Sandbrink and T. Woodside (2024). "Responsible Reporting for Frontier AI Development." In: *Proceedings of the 2024 AAAI/ACM Conference on AI, Ethics, and Society (AIES)*. URL: https://ojs.aaai.org/index.php/AIES/article/view/31735.
- Anthropic Responsible Scaling Policy; OpenAI Preparedness Framework; DeepMind Frontier Safety Framework (primary documents, most recent versions).

Discussion

METR reports that the time horizon of tasks AI agents can complete is doubling every seven months. What does this mean, and is it correct? How can we measure AI capabilities? What is construct validity, and why does it matter here?

Part III. Regulation in practice — the supply of regulation

Week 10. Tech, antitrust, and the market structure of AI

Required

- L. M. Khan (2017). "Amazon's Antitrust Paradox." In: *Yale Law Journal* 126.3, pp. 710–805. URL: https://www.yalelawjournal.org/note/amazons-antitrust-paradox.
- H. Hovenkamp (2021). "Antitrust and Platform Monopoly." In: *Yale Law Journal* 130.8, pp. 1952–2050. URL: https://www.yalelawjournal.org/article/antitrust-and-platform-monopoly.
- UK Competition and Markets Authority (2023). *AI Foundation Models: Initial Report*. Tech. rep. Competition and Markets Authority. URL: https://www.gov.uk/cma-cases/ai-foundation-models-initial-review, paired with UK Competition and Markets Authority (2024). *AI Foundation Models: Update Paper*. Tech. rep. Competition and Markets Authority. URL: https://www.gov.uk/cma-cases/ai-foundation-models-technical-update-report.

Optional

- U.S. Federal Trade Commission (Jan. 2025). *AI Partnerships and Investments: FTC Section 6(b) Study*. Tech. rep. Federal Trade Commission. URL: https://www.ftc.gov/reports/ftc-staff-report-ai-partnerships-investments-6b-study.
- C. Cunningham, F. Ederer, and S. Ma (2021). "Killer Acquisitions." In: *Journal of Political Economy* 129.3, pp. 649–702. DOI: 10.1086/712506.
- F. H. Easterbrook (1984). "The Limits of Antitrust." In: *Texas Law Review* 63.1, pp. 1–40. URL: https://chicagounbound.uchicago.edu/journal_articles/1153.
- J. M. Newman (2015). "Antitrust in Zero-Price Markets: Foundations." In: *University of Pennsylvania Law Review* 164.1, pp. 149–206. URL: https://scholarship.law.upenn.edu/penn_law_review/vol164/iss1/4/.
- A. Bradford (2023). *Digital Empires: The Global Battle to Regulate Technology*. New York: Oxford University Press. URL: https://global.oup.com/academic/product/digital-empires-9780197649268, relevant chapter.

Discussion

Should we be concerned about the market dominance of frontier AI labs? Is AI a natural monopoly?

Week 11. Who governs? Agenda-setting, GDPR, the EU AI Act, California

Required

- F. S. Berry and W. D. Berry (1990). "State Lottery Adoptions as Policy Innovations: An Event History Analysis." In: *American Political Science Review* 84.2, pp. 395–415. DOI: 10.2307/1963526.
- A. Bradford (2020). *The Brussels Effect: How the European Union Rules the World*. New York: Oxford University Press. URL: https://global.oup.com/academic/product/the-brussels-effect-9780190088583, Chapter 1 ("The Brussels Effect"); and the digital regulation chapter.
- M. Veale and F. Zuiderveen Borgesius (2021). "Demystifying the Draft EU Artificial Intelligence Act." In: *Computer Law Review International* 22.4, pp. 97–112. DOI: 10.9785/cri-2021-220402. arXiv: 2107.03721, paired with Anonymous (TBC) (2025a). "Navigating the Legislative Dilemma: Evaluating the EU AI Act's Approach to Regulating Emerging Technologies." In: *Theory and Practice of Legislation*. Author details to be confirmed before assignment. DOI: 10.1080/20508840.2025.2513177.

Optional

- F. Gilardi (2012). "Transnational Diffusion: Norms, Ideas, and Policies." In: *Handbook of International Relations*. Ed. by W. Carlsnaes, T. Risse, and B. A. Simmons. 2nd. London: SAGE Publications, pp. 453–477. URL: https://sk.sagepub.com/reference/hdbk_interntlrelations/n18.xml.
- D. Vogel (1995). *Trading Up: Consumer and Environmental Regulation in a Global Economy*. Cambridge, MA: Harvard University Press. URL: https://www.hup.harvard.edu/books/9780674900844, relevant chapter.
- Anonymous (TBC) (2025b). "Standards and the EU AI Act: Legitimacy, State of Play, and Future Challenges." In: *International Review of Law, Computers & Technology*. Author details to be confirmed before assignment. DOI: 10.1080/13600834.2025.2570966.
- Z. Li, L. Mahmoud, J. P. Quintais, et al. (2025). "Mapping the Empirical Evidence of the GDPR (In-)Effectiveness: A Systematic Review." In: *Computer Law & Security Review*. arXiv: 2310.16735.
- European Data Protection Board (2026). *Marking Ten Years of the GDPR: The Evolution of the European Data*

Protection Landscape. EDPB statement, 27 April 2026. URL: https://www.edpb.europa.eu/news/news/2026/marking-10-years-gdpr-evolution-european-data-protection-landscape_en.
- European Commission (2026). *Ten Years of the GDPR: Your Data, Your Rights*. European Commission, 22 May 2026. URL: https://commission.europa.eu/news-and-media/news/ten-years-gdpr-your-data-your-rights-2026-05-22_en.
- California SB 1047 (vetoed) and Newsom's veto message; SB 53 (signed).
- New York RAISE Act (text and signing statement).
- Encode and Calvin commentary on federal preemption of state AI law.

Discussion

Why have the EU and California been the regulators of first resort in digital and AI policy? How successful have they been? And will it last?

Week 12. National security I: drones, warfare, surveillance

Required

- P. Scharre (2018). *Army of None: Autonomous Weapons and the Future of War*. New York: W. W. Norton. URL: https://wwnorton.com/books/9780393608984, Introduction plus one substantive chapter on lethal autonomous weapon systems.
- International Committee of the Red Cross (2021). *ICRC Position on Autonomous Weapon Systems*. ICRC position paper. URL: https://www.icrc.org/en/document/icrc-position-autonomous-weapon-systems.
- Yale Law Journal Note (2021). "Introducing Independence to the Foreign Intelligence Surveillance Court." In: *Yale Law Journal*. Author of note to be confirmed before assignment. URL: https://yalelawjournal.org/note/introducing-independence-to-the-foreign-intelligence-surveillance-court.

Optional

- S. Brayne (2020). *Predict and Surveil: Data, Discretion, and the Future of Policing*. New York: Oxford University Press. URL: https://global.oup.com/academic/product/predict-and-surveil-9780190684099, relevant chapter.
- P. Scharre (2023). *Four Battlegrounds: Power in the Age of Artificial Intelligence*. New York: W. W. Norton. URL: https://wwnorton.com/books/9780393866865, relevant chapter.

Discussion

Why is LLM surveillance different? Who makes surveillance policy in the US? Might autonomous weapons be an improvement over the alternatives?

Week 13. National security II: chips, compute, China, the Qwen question

Required

- C. Miller (2022). *Chip War: The Fight for the World's Most Critical Technology*. New York: Scribner. URL: https://www.simonandschuster.com/books/Chip-War/Chris-Miller/9781982172008, Introduction plus one substantive chapter.
- G. Sastry et al. (2024). *Computing Power and the Governance of AI*. arXiv preprint. arXiv: 2402.08797.
- H. Farrell and A. L. Newman (2019). "Weaponized Interdependence: How Global Economic Networks Shape State Coercion." In: *International Security* 44.1, pp. 42–79. DOI: 10.1162/isec_a_00351.

Optional

- J. Ding (2024). *Technology and the Rise of Great Powers: How Diffusion Shapes Economic Competition*. Princeton: Princeton University Press. URL: https://press.princeton.edu/books/hardcover/9780691260334/technology-and-the-rise-of-great-powers, relevant chapter.
- Bureau of Industry and Security, October 2022 and October 2023 chip export rules (primary documents).
- H. Toner (Apr. 2026). *Testimony before the U.S. Senate Committee on the Judiciary, Subcommittee on Courts,*

Intellectual Property, and Artificial Intelligence: "Protecting Our Edge: Trade Secrets and the Global AI Arms Race". Written testimony, U.S. Senate Committee on the Judiciary. URL: https://www.judiciary.senate.gov/imo/media/doc/bd86374a-060c-a5e4-b533-d54311487456/2026-04-22_Testimony_Toner.pdf.

- H. Toner and R. Wiblin (Nov. 2025). *Helen Toner on the Geopolitics of AI in China and the Middle East*. 80,000 Hours Podcast. URL: https://80000hours.org/podcast/episodes/helen-toner-ai-policy-washington-dc/.
- M. Buyl et al. (2025). *R1dacted: Investigating Local Censorship in DeepSeek's R1 Language Model*. arXiv preprint. Authors to be confirmed; placeholder author list. arXiv: 2505.12625.
- S. Kalisz, D. Bowen, et al. (2026). *Censored LLMs as a Natural Testbed for Secret Knowledge Elicitation*. arXiv preprint. arXiv: 2603.05494.
- Authors TBC (2026). *Bilingual Bias in Large Language Models: A Taiwan Sovereignty Benchmark Study*. arXiv preprint. arXiv: 2602.06371.
- H. Waight et al. (2026). "State Media Control Influences Large Language Models." In: *Nature*. DOI: 10.1038/s41586-026-10506-7, cross-listed from Week 4.

Discussion

Is China the natural rival of the US in AI? What does this mean in practice, and what do both sides think is at stake? Should the US restrict sale of semiconductors to China?

PART IV. SYNOPSIS

Week 14. The regulatory alignment problem

Required

- N. Guha et al. (2024). "AI Regulation Has Its Own Alignment Problem: The Technical and Institutional Feasibility of Disclosure, Registration, Licensing, and Auditing." In: *George Washington Law Review* 92.6, pp. 1473–1559. URL: https://dho.stanford.edu/wp-content/uploads/AI_Regulation.pdf.

Discussion

Are existing US regulatory institutions well-prepared to meet the challenges of AI?